

Do Financial LLM Agents Discover New Alpha?

A 98-Work Audit and 50-Strategy Factor-Spanning Study

Anonymous Author(s)

Abstract

Financial LLM agents often produce strategies that look profitable, but apparent profitability is not the same as new alpha. We screen 98 distinct works, retain 69 formula-discovery or trading papers, and construct 50 good-faith U.S. strategy reconstructions for all 40 retained works with sufficient public content; the other 29 remain availability-only. Holding returns, costs, and evaluation months fixed, median annualized one-month-ahead residual return falls from 9.97% under CAPM to 3.17% under FF3, 2.69% under FF5 plus momentum, and -1.69% under the market plus 132 Jensen-Kelly-Pedersen (JKP) factors. The corresponding numbers of Holm-adjusted positive results are 6, 1, 2, and 0 out of 50. Factor overlap is direct: the median strategy's strongest absolute correlation with a JKP factor is 0.81, and 47 of 50 exceed 0.50. Among 38 mapped papers with verified full text, none reports a factor-adjusted intercept and none uses JKP132. These results support a precise interpretation: LLM agents can rediscover profitable factor-zoo ideas, yet the public evidence does not establish return information incremental to a broad known-factor library. A plausible mechanism is that finance-trained or finance-prompted agents recombine concepts prevalent in the literature while papers evaluate them against benchmarks too narrow to detect that reuse. Because the reconstructions and benchmark ladder were designed after U.S. outcomes were inspected, the estimates are descriptive spanning diagnostics, not confirmatory tests of native-agent performance.

CCS Concepts

• **Computing methodologies** → **Multi-agent systems**; • **Applied computing** → **Economics**; • **General and reference** → *Empirical studies*.

Keywords

financial agents, alpha, asset pricing, factor zoo, reproducibility, multiple testing

ACM Reference Format:

Anonymous Author(s). 2026. Do Financial LLM Agents Discover New Alpha? A 98-Work Audit and 50-Strategy Factor-Spanning Study. In *Proceedings of 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 7 pages.

1 Introduction

LLM-agent papers increasingly report profitable factors, superior trading performance, or investment alpha [14, 41, 64, 69, 93]. Many of those strategies may indeed earn positive returns relative to the papers' chosen baselines. The central empirical question, however, is whether they add return information beyond familiar asset-pricing factors. A profitable strategy that repackages low-risk, momentum, quality, or valuation exposures is not new alpha merely because an LLM generated or selected it.

Our evidence separates these claims. Across 50 reconstructions from 40 retained papers, apparent alpha remains positive against CAPM, FF3, and FF5 plus momentum, but collapses when the benchmark expands to the market plus 132 JKP characteristic factors [23]. Median annualized residual return moves from 9.97% under CAPM to -1.69% under the broad model, and the Holm-adjusted discovery count falls from six to zero. The return series themselves point to the same explanation: 47 of 50 have an absolute correlation of at least 0.50 with their nearest JKP factor, and the median nearest-factor correlation is 0.81.

This pattern is consequential because the source papers generally do not ask the spanning question. Of the 38 mapped papers whose full text we verified, 37 report no asset-pricing regression, one reports multifactor loadings without an intercept, 0 report factor-adjusted alpha, and 0 use JKP132. Consequently, a strategy can look novel under return, Sharpe, predictive, or narrow-factor comparisons while reproducing an established anomaly. The paper's claim is therefore not that all 98 screened works are "bad" or that every unavailable native agent must fail. It is that the published evidence does not substantiate the stronger claim of alpha incremental to the known factor zoo.

A plausible mechanism connects the benchmark omission to the observed overlap. Many systems prompt, retrieve, or search over financial concepts and symbolic factor expressions [41, 61, 93]. Finance text is saturated with canonical anomaly ideas, so recombination can rediscover economically meaningful signals without generating genuinely new return information. We do not observe proprietary pretraining corpora and cannot identify this mechanism causally; the benchmark ladder and correlation map establish absorption, not why a particular model selected a factor.

Contributions. We make four contributions. First, we provide a cutoff-bounded 98-work evidence waterfall that distinguishes 69 retained works, 40 reconstructed works, and 29 availability-only works without imputing failures as zero returns. Second, we evaluate all 50 retained reconstructions on one matched benchmark ladder, from CAPM through JKP132, rather than numerically back-testing only the five papers with the strongest source grounding. Third, we map every strategy return series to all 132 JKP factors and identify the factor families that absorb the claimed alpha. Fourth, we audit the benchmarks used in the mapped source papers and separate the primary factor-spanning result from a secondary 14-attempt reproducibility audit and other sensitivity analyses.

One-sentence takeaway. Financial LLM agents often recover strategies that look like alpha against narrow benchmarks, but the reconstructed returns mostly repackage known factor-zoo exposures and provide no familywise evidence of incremental alpha beyond JKP132.

2 Evidence Design

2.1 Corpus, denominators, and evidence routes

The screen covers database inception through 2 August 2026. Three query families combine LLM/agent terms with alpha or strategy discovery; trading, portfolios, returns, or backtests; and benchmark, audit, artifact, or leakage terms. Sources are arXiv’s relevant CS and q-fin categories, ACL Anthology, OpenReview, ACM DL, and SSRN, plus citation chasing from AlphaBench, AlphaQT-Bench, FINSABER, and the Agentic Trading survey. Crossref/OpenAlex resolve duplicates; GitHub locates only record-linked artifacts. Raw hit counts and rankings were not preserved, so this is a cutoff-bounded systematic screen rather than a complete universe.

The registry has 103 lineages, 104 record URLs, and 98 works after six duplicate manifestations are collapsed. All 98 works are cited. Frozen rules retain 30 formula-discovery and 39 trading works [4, 9, 14, 25, 28, 41, 46, 49, 57, 61, 64, 67, 69, 72, 83, 91] [3, 20, 29, 39, 42, 45, 60, 65, 71, 73, 85, 87, 92, 93] [15, 40, 68, 76, 89, 90, 95] [1, 2, 7, 17, 18, 27, 36–38, 43, 44, 47, 51, 54, 62, 63, 66, 74, 75, 78, 79, 97] [5, 6, 11, 16, 22, 24, 33, 35, 56, 70] and exclude 29 benchmarks, audits, community tools, or mechanistic comparators with recorded reasons [10, 58, 59, 80, 88, 94] [8, 12, 13, 31, 32, 52, 82, 86] [19, 26, 30, 34, 48, 50, 53, 55, 77, 81, 84, 96, 98–100]. Thus 98/69/40/50 denote screened works, retained works, reconstructed works, and reconstructed strategies; they are not competing replication counts.

We construct 50 mappings for 40 retained works. Five works support 13 source-grounded component tests, while 35 works support 37 favorable in-spirit reconstructions. The other 29 retained works remain availability-only because no frozen common-panel mapping is supportable; we neither invent a proxy nor assign zero alpha. One of the 50 strategies adapts a released QuantEvolver seed expression, 12 reconstruct source-grounded components, and 37 translate higher-level narratives. Every estimate inherits that provenance: a reconstruction tests the historical economic content we implemented, not the unavailable native agent.

2.2 Common U.S. task

The monthly U.S. panel runs from August 2001 through December 2024. At formation month t , the universe $\mathcal{U}_t$ contains the 1,000 largest positive-market-equity common stocks, and mapping i produces point-in-time score $s_{ijt} = g_i(x_{jt})$. For the standard value-weighted long–short mapping, with top and bottom deciles H_{it} and L_{it} and market equity $m_{jt} > 0$, target weights are

$$w_{ijt} = \begin{cases} m_{jt} / \sum_{k \in H_{it}} m_{kt}, & j \in H_{it}, \\ -m_{jt} / \sum_{k \in L_{it}} m_{kt}, & j \in L_{it}, \\ 0, & \text{otherwise.} \end{cases} \quad (1)$$

Labeled long-only and sparse top- k mappings preserve their stated rules. Formation never conditions on next-month return availability. Missing held returns receive zero in the primary analysis without renormalizing surviving weights; an adverse unit-move stress is reported separately.

Turnover restores target exposure from drifted pre-trade weights. With total security return $R_{j,t}^{\text{tot}}$ and held excess return $R_{j,t+1}^{e,*}$,

$$\begin{aligned} \tilde{w}_{ij,t-} &= \frac{w_{ij,t-1}(1 + R_{j,t}^{\text{tot}})}{1 + \sum_k w_{ik,t-1} R_{k,t}^{\text{tot}}}, \quad \tau_{it} = \sum_j |w_{ijt} - \tilde{w}_{ij,t-}|, \\ r_{i,t+1}^{\text{net}}(c) &= \sum_j w_{ijt} R_{j,t+1}^{e,*} - c\tau_{it}, \quad c \in \{0, 5, 10, 25, 50\} \times 10^{-4}. \end{aligned} \quad (2)$$

The primary one-way rate is 10 bp. This linear cost model omits borrow, financing, nonlinear impact, and capacity; gross benchmark factors make the comparison conservative.

2.3 Matched benchmark ladder and factor overlap

The headline analysis restricts attention to all 50 strategies associated with the 40 reconstructed retained works and fixes their net returns, 126 evaluation months (August 2011–January 2022), and inference across four benchmarks. For benchmark b , slopes use a rolling 120-month window. The broad model contains the market and 132 JKP characteristic returns; the six FF5-plus-momentum analogues remain unpenalized, and ridge penalizes the other 127 slopes. The penalty $\lambda \in \{0.1, 1, 10, 100, 1000\}$ minimizes mean-squared error over the last 24 months of each training window before refitting on all 120 months. CAPM, FF3, and FF5 plus momentum use unpenalized rolling OLS. For each test month,

$$\hat{\theta}_{it}^{(b)}(\lambda) = \arg \min_{\theta=(a,\beta)} \left\{ \sum_{s=t-120}^{t-1} (r_{is} - a - \beta^\top f_s^{(b)})^2 + \lambda \|P_b \beta\|_2^2 \right\}, \quad (3)$$

$$e_{it}^{(b)} = r_{it} - \hat{\beta}_{it}^{(b)\top} f_t^{(b)}, \quad \hat{\alpha}_i^{(b)} = 12 \bar{e}_i^{(b)}, \quad (4)$$

where P_b exempts the six base returns from penalization. Omitting the fitted intercept from $e_{it}^{(b)}$ makes its mean the factor-unexplained return. HAC statistics use $L = 4$; Holm and a 5,000-replication moving-block maximum- $|t|$ procedure control the 50-strategy family conditional on the mappings.

To identify familiar economic content directly, we correlate each strategy with every JKP factor over the 246 common months:

$$\rho_{ik} = \text{Corr}(r_i, f_k^{\text{JKP}}), \quad k_i^* = \arg \max_{1 \leq k \leq 132} |\rho_{ik}|. \quad (5)$$

This nearest-factor statistic is descriptive. Correlation does not identify the agent’s inputs or prove that a source paper copied a factor; it tests whether the reconstructed return path resembles a known factor return.

3 Results

3.1 Apparent alpha is absorbed by known factors

Figure 1 states the paper’s main result. Under CAPM, 44 of 50 strategies have positive estimated alpha, 28 are nominally positive, and 6 survive Holm; median annualized alpha is 9.97%. FF3 reduces the median to 3.17% and leaves 1 Holm result. FF5 plus momentum yields median 2.69% and 2 Holm results. Once the benchmark

Table 1: The five papers underlying the 13 source-grounded component tests. “Source-grounded” applies to the named component, not the full native agent or original evaluation. The other 37 retained strategies are explicitly labeled in-spirit reconstructions.

Source (<i>N</i> mappings)	Original idea and task	Component implemented in our common task
EFS [49] (3)	LLM evolutionary factor search for sparse top- <i>m</i> portfolios on Fama–French and US50/HSI45/CSI300 datasets.	Momentum/low-volatility/breakout score; reversal/low-noise score; and a sparse equal-weight top-five variant.
QuantEvolver [93] (1)	Reinforcement fine-tuning generates executable factor-DSL expressions across market benchmarks.	Released risk-adjusted-momentum seed $\text{ret}_{12,1}/ \text{rvol}_{252} $; portfolio formation remains ours.
Alpha-Jungle [61] (2)	LLM-guided MCTS searches symbolic stock factors using price, volume, and volatility operators.	Price/volume/momentum and trend-after-volatility-compression scores based on documented examples.
FAMA [41] (1)	Neural-symbolic factor mining evaluated by RankIC/RankICIR and an S&P 500 simulation.	Interpretable value–momentum–profitability score motivated by the paper’s examples.
GuruAgents [27] (6)	Prompt-guided agents emulate named investor philosophies in a Nasdaq-100 backtest.	Graham, Buffett, Greenblatt, Piotroski, and Altman screens plus their ensemble.

includes JKP132, the median becomes -1.69%: only 17 estimates remain positive, one is nominally positive, and none survives Holm or maximum- $|t|$ adjustment. The median attenuation from FF5 plus momentum to JKP132 is 4.95% per year.

The result is not driven only by the 37 in-spirit translations. Among the 13 released-expression or source-grounded component strategies, 10 have positive FF3 alpha and 11 have positive FF5-plus-momentum alpha, but only 4 remain positive under JKP132. The benchmark expansion therefore changes the economic conclusion even for the highest-fidelity subset. It also clarifies the role of the five papers in Table 1: they define an attribution tier, not the numerical backtest denominator.

3.2 Which known factors absorb the returns?

The correlation map supplies the missing economic interpretation. Across the 50 strategies, the median $|\rho_{iK_i^*}|$ is 0.81; 47 have a nearest-factor absolute correlation of at least 0.50. Although 21 different JKP characteristics appear as a closest match, Table 2 shows concentration in low-risk, 52-week-high, volatility, quality, momentum, and valuation families. Betting-against-beta alone is the nearest analogue for 12 strategies and 52-week-high proximity for eight.

The signs are economically coherent. Risk-aware and stable-quality systems often align with low-beta or low-volatility returns. Their correlations with positively oriented volatility factors can therefore be negative. Trend and breakout strategies align positively with 52-week-high and momentum factors, while fundamental agents align with quality or valuation. JKP132 does not merely add statistical flexibility: it spans the same economic motifs that dominate the reconstructed agent strategies.

The source benchmark audit explains why papers can miss this absorption. Among 38 verified mapped papers, 37 contain no asset-pricing factor regression and 1 reports loadings without a factor-adjusted intercept; none reports factor-adjusted alpha or uses JKP132. Source evaluations instead emphasize predictive scores, cumulative or excess returns, Sharpe ratios, buy-and-hold comparisons, agent baselines, or factor loadings. Those outcomes can establish usefulness relative to a chosen task, but they cannot establish incremental alpha beyond the factor zoo.

3.3 Fidelity, costs, and secondary diagnostics

The longer 2001–2024 six-factor exercise provides an implementation sensitivity, not the headline matched ladder. Among the 13 source-grounded component tests, median annualized six-factor

alpha is 0.88%; one is nominally positive and none survives Holm within that subset. Across the broader 62-row ledger, which adds 12 mappings from excluded diagnostic sources, median alpha at 10 bp is -0.05% and one researcher-authored AlphaAgent-inspired quality proxy survives conditional Holm. That result is not agent output, disappears under some fixed HAC choices, and has a simultaneous lower bound of only 0.48%.

Trading costs weaken rather than create the result. Median six-factor alpha falls from 2.15% gross to 1.01% at 5 bp and -0.05% at 10 bp; positive counts across 0/5/10/25/50 bp are 46/42/30/18/10. An adverse missing-return stress leaves zero nominal or Holm positives. A separate 133-factor diagnostic over the 62-row ledger yields one nominal result but zero Holm or maximum- $|t|$ results.

The international extension is forensic rather than performance evidence. Thirty-five paths cross nonpositive NAV; security-level reconstruction shows that every event is dominated by one extreme positive return held short. Without independent validation of those returns and corporate actions, we retain the event ledger but exclude international alpha from the paper’s conclusion.

4 Interpretation and Limits

The evidence supports a strong but bounded statement. LLM-agent strategies can produce economically meaningful returns and may automate factor discovery, formula search, or portfolio construction. What the public results do not establish is *new* alpha: against JKP132, the reconstructed family has negative median residual return and no multiplicity-adjusted positive result. The direct correlation evidence links that attenuation to familiar factor families rather than treating a larger benchmark as an unexplained statistical black box.

The proposed mechanism remains a hypothesis. Financial LLM systems operate on pretrained knowledge, prompts, retrieved text, hand-specified operators, or factor DSLs that can foreground established anomalies. If evaluation then uses buy-and-hold, Sharpe, predictive metrics, or FF3/FF5, rediscovery can be reported as alpha because the benchmark omits the relevant exposure. Our audit verifies the second link—broad factor adjustment is absent—and the return correlations verify factor overlap. It does not observe proprietary training examples or causally attribute any strategy to memorization.

Three limitations prevent a stronger agent-level claim. First, 37 of 50 strategies are favorable in-spirit reconstructions, 12 are source-grounded component tests, and only one adapts released code; none is a native-agent return replication. Second, mappings and




Figure 1: Matched benchmark ladder for all 50 retained strategy reconstructions from 40 papers, net of 10 bp one-way costs. Candidate returns and the 126 evaluation months are identical across models. CAPM, FF3, and FF5 plus momentum use rolling OLS; the broad specification leaves the six base factors unpenalized and ridge-controls the remaining JKP factors. Annotations report positive estimates, positive two-sided HAC tests at 5%, and positive Holm-adjusted tests. Because the mappings and ladder are post-hoc, “one-month-ahead” describes factor fitting, not prospective agent discovery.

Table 2: Most frequent JKP132 analogues of the 50 retained strategy reconstructions. The table reports the JKP factor with the largest absolute return correlation for each strategy, not a factor label claimed verbatim by the source paper.

Closest JKP analogue	Strategies	Median ρ	Median $ \rho $
Betting-against-beta	12	-0.84	0.84
52-week-high proximity	8	0.85	0.85
Realized volatility	3	-0.82	0.82
Idiosyncratic volatility	3	-0.78	0.78
Quality-minus-junk safety	3	0.77	0.77
12-1 momentum	3	0.71	0.71
Downside beta	2	-0.19	0.70
EBITDA/enterprise value	2	0.69	0.69
3-1 momentum	2	0.66	0.66
Average five highest daily returns	1	-0.84	0.84

Correlations use 246 common monthly returns from August 2001 through January 2022. Negative ρ denotes exposure opposite to the published JKP factor orientation. These ten analogues cover 39 strategies; eleven additional JKP factors are the closest match for one strategy each.

the benchmark ladder were frozen after U.S. outcomes had been inspected and had no independent second coder. Holm and bootstrap procedures cover sampling multiplicity across realized formulas, not mapping discretion, prior factor mining, or publication selection [21]. Third, common-task portfolios omit language, memory, retrieval, debate, native timing, dynamic allocation, non-equities,

and execution. A reconstruction failure cannot prove native-agent failure, and reconstruction success cannot be attributed to an agent.

5 Reproducibility and Secondary Code Audit

The secondary repository audit targets 14 implementations: eight retained systems and six excluded diagnostic works. It reproduces

zero native agents. One released QuantEvolver seed expression supports a testable code-backed adaptation with annualized six-factor alpha of 0.33% ($t = 0.65$); the other 13 attempts have recorded output, input/adaptor, task, licensing, or executability blockers and do not enter a performance denominator. Keeping this audit secondary avoids conflating repository availability with the 40-paper/50-strategy spanning analysis.

The artifact ledger contains the search protocol and 98 dispositions; the 50 strategy definitions and provenance fields; the source-paper benchmark coding; all 6,600 strategy–JKP correlations; the four-model benchmark panel; cost, missing-return, and multiplicity outputs; and the 14 code-attempt records. Deterministic builders verify hashes and regenerate tables and figures. Tests cover timing, weights, drift, costs, factor alignment, rolling estimation, denominators, correlation completeness, and claim values.

Responsible use and AI assistance. No orders were transmitted. Results are retrospective research evidence, not investment advice. Generative AI assisted literature organization, code drafting, diagnostic review, and prose editing; the human author is responsible for source selection, mappings, protocol choices, statistical specifications, citation verification, and all reported claims.

6 Conclusion

Financial LLM-agent papers often identify profitable strategies; our result is not that those returns are fictitious. The result is that profitability against a narrow benchmark is being mistaken for new alpha. Across all 50 retained reconstructions from 40 papers, median annualized residual return falls from 9.97% under CAPM to -1.69% under the market plus JKP132, and the Holm-adjusted discovery count falls to zero. The median strategy has 0.81 absolute correlation with its closest JKP factor, with the strongest concentration in low-risk, 52-week-high, volatility, quality, momentum, and valuation families.

The evidence supports a coherent story: finance-oriented agents recombine ideas already abundant in the factor literature, while source papers rarely use a benchmark broad enough to distinguish rediscovery from innovation. This mechanism is plausible rather than causally identified, but the empirical implication is clear. Claims of LLM-generated alpha should report broad factor spanning, factor overlap, multiplicity correction, costs, and dated reproducible outputs before treating agent-produced returns as incremental investment information.

References

- [1] Blake Almon, Matthew Caliboso, Alex Kim, Rohan Dutta, Rohan Raman, Mithil Srungarapu, Vasu Sharma, and Kevin Zhu. 2025. RAPTOR: Reasoned Agentic Portfolio Trading with Orchestrated Rebalancing. In *NeurIPS 2025 Workshop on New Opportunities for Research and Applications in Finance*. <https://openreview.net/forum?id=ziuTkKhgT0>
- [2] Mostapha Benhenda. 2025. FinRL-DeepSeek: LLM-Infused Risk-Sensitive Reinforcement Learning for Trading Agents. arXiv:2502.07393 [q-fin.TR] <https://arxiv.org/abs/2502.07393>
- [3] Mehmet Caner, Agostino Capponi, Nathan Sun, and Jonathan Y. Tan. 2026. Designing Agentic AI-Based Screening for Portfolio Investment. arXiv:2603.23300 [q-fin.PM] <https://arxiv.org/abs/2603.23300>
- [4] Lang Cao. 2025. Chain-of-Alpha: Unleashing the Power of Large Language Models for Alpha Mining in Quantitative Trading. arXiv:2508.06312 [cs.CE] <https://arxiv.org/abs/2508.06312>
- [5] Jiaxiang Chen, Aotian Luo, Zhouyi Zheng, Weiwei Huang, Chi Zhang, and Zenglin Xu. 2026. MetaPS: Adaptive Programmatic Strategy Selection for Market Agents. arXiv:2606.22385 [cs.AI] <https://arxiv.org/abs/2606.22385>
- [6] Xiwen Chen, Wenhui Zhu, Songzhu Zheng, Kashif Rasul, Yueyue Deng, and Huayu Li. 2026. SHARP: A Self-Evolving Human-Auditable Rubric Policy for Financial Trading Agents. arXiv:2605.06822 [cs.LG] <https://arxiv.org/abs/2605.06822>
- [7] Yang Chen, Yuchen Cao, Jacky Keung, Leilei Gan, Kun Kuang, Yueheng Jiang, Zhaozhao Ma, Jianping Zhu, Fei Wu, and Jinpeng Li. 2025. MM-ARC: Multimodal Adaptive Routing of Capital with Robustness-Audited Strategy Pools. arXiv:2509.05080 [q-fin.TR] <https://arxiv.org/abs/2509.05080>
- [8] Yanxu Chen, Zijun Yao, Yantao Liu, Amy Xin, Jin Ye, Jianing Yu, Lei Hou, and Juanzi Li. 2025. StockBench: Can LLM Agents Trade Stocks Profitably In Real-world Markets? arXiv:2510.02209 [cs.LG] <https://arxiv.org/abs/2510.02209>
- [9] Junyan Cheng and Peter Chin. 2024. Empirical Asset Pricing with Large Language Model Agents. arXiv:2409.17266 [cs.AI] <https://arxiv.org/abs/2409.17266>
- [10] Can Cui, Wei Wang, Meihui Zhang, Gang Chen, Zhaojing Luo, and Beng Chin Ooi. 2021. AlphaEvolve: A Learning Framework to Discover Novel Alphas in Quantitative Investment. arXiv:2103.16196 [cs.AI] <https://arxiv.org/abs/2103.16196>
- [11] Zheyang Deng, Weixiang Yan, Changlong Yu, and Jiahu Wang. 2026. AlphaQuant: An End-to-End Tool-Augmented Agentic Reinforcement Learning Framework for Stock Trading. In *Findings of the Association for Computational Linguistics: ACL 2026*. <https://aclanthology.org/2026.findings-acl.456/>
- [12] Momoza Dimodugno and Mehdi Mammadov. 2025. AICrypto-Assistant: A Multi-Agent LLM Platform for Democratizing Crypto-Asset Analysis. In *International Conference on Open Finance Proceedings*. doi:10.36287/setsci.24.2.017
- [13] Hongjun Ding, Binqi Chen, Jinsheng Huang, Taian Guo, Zhengyang Mao, Guoyi Shao, Lutong Zou, Luchen Liu, and Ming Zhang. 2025. AlphaEval: A Comprehensive and Efficient Evaluation Framework for Formula Alpha Mining. arXiv:2508.13174 [cs.AI] <https://arxiv.org/abs/2508.13174>
- [14] Yitong Duan, chuheng zhang, and Jian Li. 2025. FactorMAD: A Multi-Agent Debate Framework Based on Large Language Models for Interpretable Stock Alpha Factor Mining. In *Proceedings of the 6th ACM International Conference on AI in Finance*. doi:10.1145/3768292.3770377
- [15] Sorourasadat Fatemi and Yuheng Hu. 2024. FinVision: A Multi-Agent Framework for Stock Market Prediction. arXiv:2411.08899 [q-fin.TR] <https://arxiv.org/abs/2411.08899>
- [16] George Fatouros and Kostas Metaxas. 2026. Signal or Noise in Multi-Agent LLM-based Stock Recommendations? arXiv:2604.17327 [q-fin.PM] <https://arxiv.org/abs/2604.17327>
- [17] George Fatouros, Kostas Metaxas, John Soldatos, and Manos Karathanassis. 2025. MarketSenseAI 2.0: Enhancing Stock Analysis through LLM Agents. arXiv:2502.00415 [q-fin.CP] <https://arxiv.org/abs/2502.00415>
- [18] Taian Guo, Haiyang Shen, JinSheng Huang, Zhengyang Mao, Junyu Luo, Binqi Chen, Zhuoru Chen, Luchen Liu, Bingyu Xia, Xuhui Liu, Yun Ma, and Ming Zhang. 2025. MASS: Multi-agent simulation scaling for portfolio construction. arXiv:2505.10278 [cs.AI] <https://arxiv.org/abs/2505.10278>
- [19] Taian Guo, Haiyang Shen, Junyu Luo, Binqi Chen, Hongjun Ding, Jinsheng Huang, Luchen Liu, Yun Ma, and Ming Zhang. 2026. AlphaPROBE: Alpha Mining via Principled Retrieval and On-graph biased evolution. arXiv:2602.11917 [cs.AI] <https://arxiv.org/abs/2602.11917>
- [20] Jun Han, Shuo Zhang, Wei Li, Yifan Dong, Tu Hu, Yumo Zhu, Xiaomin Yu, Xin Guo, Zhaowei Liu, Kunyi Wang, Jingping Liu, Tianyi Jiang, Ruichuan An, Sen Hu, Zhi Yang, Ronghao Che, and Huacan Wang. 2026. QuantaAlpha: An Evolutionary Framework for LLM-Driven Alpha Mining. arXiv:2602.07085 [q-fin.ST] <https://arxiv.org/abs/2602.07085>
- [21] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *The Review of Financial Studies* 29, 1 (2016), 5–68. doi:10.1093/rfs/hhv059
- [22] Wenliang Huang and Zengyi Yu. 2026. Strat-LLM: Stratified Strategy Alignment for LLM-based Stock Trading with Real-time Multi-Source Signals. arXiv:2605.06024 [cs.AI] <https://arxiv.org/abs/2605.06024>
- [23] Theis Ingerslev Jensen, Bryan T. Kelly, and Lasse Heje Pedersen. 2023. Is There a Replication Crisis in Finance? *The Journal of Finance* 78, 5 (2023), 2465–2518. doi:10.1111/jofi.13249
- [24] Joohyoung Jeon and Hongchul Lee. 2026. Can Blindfolded LLMs Still Trade? An Anonymization-First Framework for Portfolio Optimization. arXiv:2603.17692 [cs.LG] <https://arxiv.org/abs/2603.17692>
- [25] Zuoyou Jiang, Li Zhao, Rui Sun, Ruohan Sun, Zhongjian Li, Jing Li, Daxin Jiang, Zuo Bai, and Cheng Hua. 2025. Alpha-R1: Alpha Screening with LLM Reasoning via Reinforcement Learning. arXiv:2512.23515 [q-fin.TR] <https://arxiv.org/abs/2512.23515>
- [26] Alexey Khoroshilov, Alexey Chernysh, Orkhan Ekhtibarov, Nini Kamkia, and Dmitry Zmitrovich. 2026. QuantCode-Bench: A Benchmark for Evaluating the Ability of Large Language Models to Generate Executable Algorithmic Trading Strategies. arXiv:2604.15151 [cs.CL] <https://arxiv.org/abs/2604.15151>
- [27] Yejin Kim, Youngbin Lee, Juhyeong Kim, and Yongjae Lee. 2025. GuruAgents: Emulating Wise Investors with Prompt-Guided LLM Agents. arXiv:2510.01664 [cs.AI] <https://arxiv.org/abs/2510.01664>

- [28] Zhizhuo Kou, Holam Yu, Junyu Luo, Jingshu Peng, Xujia Li, Chengzhong Liu, Juntao Dai, Lei Chen, Sirui Han, and Yike Guo. 2025. Automate Strategy Finding with LLM in Quant Investment. In *Findings of the Association for Computational Linguistics: EMNLP 2025*. <https://aclanthology.org/2025.findings-emnlp.1005/>
- [29] Yuri Kvasiuk, Tianyi Li, Owen Colegrove, and Moritz Münchmeyer. 2026. MadEvol: Evolutionary Optimization of Trading Systems with Large Language Models. arXiv:2605.23007 [q-fin.TR] <https://arxiv.org/abs/2605.23007>
- [30] Changlun Li, Peixian Ma, Qiqi Duan, Zhenyu Lin, and Peineng Wu. 2026. NextFund: A Unified Performance Tracking Platform for Agentic Portfolio Management. arXiv:2607.11141 [cs.AI] <https://arxiv.org/abs/2607.11141>
- [31] Changlun Li, Yao Shi, Chen Wang, Qiqi Duan, Runke Ruan, Weijie Huang, Haonan Long, Lijun Huang, Nan Tang, and Yuyu Luo. 2025. Time Travel is Cheating: Going Live with DeepFund for Real-Time Fund Investment Benchmarking. arXiv:2505.11065 [cs.CE] <https://arxiv.org/abs/2505.11065>
- [32] Haoqiang Li, Yupeng Cao, Yangyang Yu, Shashidhar Reddy Javaji, Zhiyang Deng, Yueru He, Yuechen Jiang, Zining Zhu, K.P. Subbalakshmi, Jimin Huang, Lingfei Qian, Xueqing Peng, Jordan W. Suchow, and Qianqian Xie. 2025. INVESTOR-BENCH: A Benchmark for Financial Decision-Making Tasks with LLM-based Agent. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. <https://aclanthology.org/2025.acl-long.126/>
- [33] Minghan Li, Rachel Gonsalves, Weiyue Li, Sunghoon Yoon, and Mengyu Wang. 2026. TrustTrade: Human-Inspired Selective Consensus Reduces Decision Uncertainty in LLM Trading Agents. arXiv:2603.22567 [cs.CE] <https://arxiv.org/abs/2603.22567>
- [34] Weixian Waylon Li, Hyeonjun Kim, Mihai Cucuringu, and Tiejun Ma. 2026. Can LLM-based Financial Investing Strategies Outperform the Market in Long Run?. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.1*. doi:10.1145/3770854.3785702
- [35] Xiang Li, Zikai Wei, Yiyao Qi, Wanyun Zhou, Xiang Liu, Penglei Sun, Jian Guo, Yongqi Zhang, and Xiaowen Chu. 2026. Janus-Q: End-to-End Event-Driven Trading via Hierarchical-Gated Reward Modeling. arXiv:2602.19919 [cs.CL] <https://arxiv.org/abs/2602.19919>
- [36] Xiangyu Li, Yawen Zeng, Xiaofen Xing, Jin Xu, and Xiangmin Xu. 2025. HedgeAgents: A Balanced-aware Multi-agent Financial Trading System. arXiv:2502.13165 [cs.MA] <https://arxiv.org/abs/2502.13165>
- [37] Xiangyu Li, Yawen Zeng, Xiaofen Xing, Jin Xu, and Xiangmin Xu. 2025. Profit Mirage: Revisiting Information Leakage in LLM-based Financial Agents. arXiv:2510.07920 [cs.AI] <https://arxiv.org/abs/2510.07920>
- [38] Xiangyu Li, Yawen Zeng, Xiaofen Xing, Jin Xu, and Xiangmin Xu. 2025. QuantAgents: Towards Multi-agent Financial System via Simulated Trading. arXiv:2510.04643 [cs.AI] <https://arxiv.org/abs/2510.04643>
- [39] Yuqi Li, Siyuan Liu, and Bingjun Liu. 2026. AI Trading's Alpha Singularity: Emergent Market Reasoning through Agent-to-Agent Self-Evolution. arXiv:2606.29194 [cs.AI] <https://arxiv.org/abs/2606.29194>
- [40] Yuan Li, Bingqiao Luo, Qian Wang, Nuo Chen, Xu Liu, and Bingsheng He. 2024. CryptoTrade: A Reflective LLM-based Agent to Guide Zero-shot Cryptocurrency Trading. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. <https://aclanthology.org/2024.emnlp-main.63/>
- [41] Zhiwei Li, Ran Song, Caihong Sun, Wei Xu, Zhengtao Yu, and Ji-Rong Wen. 2024. Can Large Language Models Mine Interpretable Financial Factors More Effectively? A Neural-Symbolic Factor Mining Agent Model. In *Findings of the Association for Computational Linguistics: ACL 2024*. <https://aclanthology.org/2024.findings-acl.233/>
- [42] Qinhong Lin, Ruitao Feng, Yinglun Feng, Zhenxin Huang, Yukun Chen, Zhongliang Yang, Linna Zhou, Binjie Fei, Jiaqi Liu, and Yu Li. 2026. FactorEngine: A Program-level Knowledge-Infused Factor Mining Framework for Quantitative Investment. arXiv:2603.16365 [cs.AI] <https://arxiv.org/abs/2603.16365>
- [43] Bijia Liu and Ronghao Dang. 2025. FinPos: A Position-Aware Trading Agent System for Real Financial Markets. arXiv:2510.27251 [cs.MA] <https://arxiv.org/abs/2510.27251>
- [44] Bijia Liu and Ronghao Dang. 2025. FINRS: A Risk-Sensitive Trading Framework for Real Financial Markets. arXiv:2511.12599 [cs.MA] <https://arxiv.org/abs/2511.12599>
- [45] Fengyuan Liu, Yuchen Fu, Yuqi Wang, and Qi Liu. 2026. XALPHA: A Memory-Driven AI Quant Researcher for Hypothesis-to-Code Alpha Discovery. arXiv:2607.08332 [cs.CL] <https://arxiv.org/abs/2607.08332>
- [46] Fengyuan Liu, Yi Huang, Sichun Luo, Yuqi Wang, Yazheng Yang, Xinye Li, Zefa Hu, Junlan Feng, and Qi Liu. 2025. Cognitive Alpha Mining via LLM-Driven Code-Based Evolution. arXiv:2511.18850 [cs.CL] <https://arxiv.org/abs/2511.18850>
- [47] Chen-Che Lu, Yun-Cheng Chou, and Teng-Ruei Chen. 2025. P1GPT: a multi-agent LLM workflow module for multi-modal financial information analysis. arXiv:2510.23032 [cs.CE] <https://arxiv.org/abs/2510.23032>
- [48] Haochen Luo, Ho Tin Ko, Jiaodong Chen, David Sun, Yuan Zhang, and Chen Liu. 2026. AlphaBench: Benchmarking Large Language Models in Formulaic Alpha Factor Mining. In *The Fourteenth International Conference on Learning Representations*. <https://alphanbench.co/>
- [49] Haochen Luo, Yuan Zhang, and Chen Liu. 2025. EFS: Evolutionary Factor Searching for Sparse Portfolio Optimization Using Large Language Models. arXiv:2507.17211 [q-fin.PM] <https://arxiv.org/abs/2507.17211>
- [50] Sichun Luo, Yi Huang, Shichang Meng, Fengyuan Liu, Mukai Li, Qinghua Yao, Zefa Hu, Junlan Feng, and Qi Liu. 2026. AlphaQT-Bench: Diagnosing the Gap between Financial Code Generation and Quantitative Reasoning in LLMs. In *Findings of the Association for Computational Linguistics: ACL 2026*. <https://aclanthology.org/2026.findings-acl.138/>
- [51] Yichen Luo, Yebo Feng, Jiahua Xu, Paolo Tasca, and Yang Liu. 2025. LLM-Powered Multi-Agent System for Automated Crypto Portfolio Management. arXiv:2501.00826 [q-fin.TR] <https://arxiv.org/abs/2501.00826>
- [52] Tianmi Ma, Jiawei Du, Wenxin Huang, Wenjie Wang, Liang Xie, Xian Zhong, and Joey Tianyi Zhou. 2025. Agent Trading Arena: A Study on Numerical Understanding in LLM-Based Agents. In *Findings of the Association for Computational Linguistics: EMNLP 2025*. <https://aclanthology.org/2025.findings-emnlp.294/>
- [53] Andrei Neagu, Eeham Khan, and Leila Kosseim. 2026. CLaC@FinMMEval 2026 Task 3: Sentiment-Augmented Deep Reinforcement Learning for Active Trading – An Alpha-Reward Approach. arXiv:2607.16028 [cs.LG] <https://arxiv.org/abs/2607.16028>
- [54] Charidimos Papadakis, Angeliki Dimitriou, Giorgos Filandrianos, Maria Lymperaiou, Konstantinos Thomas, and Giorgos Stamou. 2025. ATLAS: Adaptive Trading with LLM Agents Through Dynamic Prompt Optimization and Multi-Agent Coordination. arXiv:2510.15949 [q-fin.TR] <https://arxiv.org/abs/2510.15949>
- [55] Lingfei Qian, Xueqing Peng, Yan Wang, Vincent Jim Zhang, Huan He, Hanley Smith, Yi Han, Yueru He, Haoqiang Li, Yupeng Cao, Yangyang Yu, Alejandro Lopez-Lira, Peng Lu, Jian-Yun Nie, Guojun Xiong, Jimin Huang, and Sophia Ananiadou. 2026. When Agents Trade: Live Multi-Market Trading Arena for LLM Agents. In *Proceedings of the ACM Web Conference 2026*. doi:10.1145/3774904.3792821
- [56] Mohotarema Rashid, Lingzi Hong, Junhua Ding, and K. S. M. Tozammel Hossain. 2026. Fin-Analyst at FinMMEval 2026 Task 3: A Live Hybrid Trading Agent with LLM Specialists and Rule-Based Signals. arXiv:2607.12233 [cs.CL] <https://arxiv.org/abs/2607.12233>
- [57] Junji Ren, Junjie Zhao, Shengcai Liu, and Peng Yang. 2025. From Flat to Hierarchical: Evolving Tree-structured Thoughts for Fine-grained Alpha Mining. arXiv:2508.16334 [cs.CE] <https://arxiv.org/abs/2508.16334>
- [58] Tao Ren, Ruihan Zhou, Jinyang Jiang, Jiafeng Liang, Qinghao Wang, and Yijie Peng. 2024. RiskMiner: Discovering Formulaic Alphas via Risk Seeking Monte Carlo Tree Search. arXiv:2402.07080 [q-fin.CP] <https://arxiv.org/abs/2402.07080>
- [59] Hao Shi, Weili Song, Xinting Zhang, Jiahe Shi, Cuicui Luo, Xiang Ao, Hamid Arian, and Luis Seco. 2024. AlphaForge: A Framework to Mine and Dynamically Combine Formulaic Alpha Factors. arXiv:2406.18394 [q-fin.CP] <https://arxiv.org/abs/2406.18394>
- [60] Runze Shi, Shengyu Yan, Yuecheng Cai, and Chengxi Lv. 2026. Hubble: An LLM-Driven Agentic Framework for Safe, Diverse, and Reproducible Alpha Factor Discovery. arXiv:2604.09601 [cs.AI] <https://arxiv.org/abs/2604.09601>
- [61] Yu Shi, Yitong Duan, and Jian Li. 2025. Navigating the Alpha Jungle: An LLM-Powered MCTS Framework for Formulaic Factor Mining. arXiv:2505.11122 [cs.AI] <https://arxiv.org/abs/2505.11122>
- [62] Zifan Song, Kaitao Song, Guosheng Hu, Ding Qi, Junyao Gao, Xiaohua Wang, Dongsheng Li, and Cairong Zhao. 2025. Trade in Minutes! Rationality-Driven Agentic System for Quantitative Financial Trading. arXiv:2510.04787 [cs.MA] <https://arxiv.org/abs/2510.04787>
- [63] Rui Sun, Li Zhao, Zuoyou Jiang, Bo Yang, Yuxiao Bai, Mengting Chen, Jing Li, and Zuo Bai. 2025. ContestTrade: A Multi-Agent Trading System Based on Internal Contest Mechanism. arXiv:2508.00554 [q-fin.TR] <https://arxiv.org/abs/2508.00554>
- [64] Ziyi Tang, Zechuan Chen, Jiarui Yang, Jiayao Mai, Yongsun Zheng, Keze Wang, Jinrui Chen, and Liang Lin. 2025. AlphaAgent: LLM-Driven Alpha Mining with Regularized Exploration to Counteract Alpha Decay. arXiv:2502.16789 [cs.CE] <https://arxiv.org/abs/2502.16789>
- [65] Ziyi Tang, Xuexiong Yin, Weixing Chen, Zechuan Chen, Yongsun Zheng, Wenxuan Ye, Keze Wang, and Liang Lin. 2026. AlphaAgentEvo: Evolution-Oriented Alpha Mining via Self-Evolving Agentic Reinforcement Learning. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=INmZrawUMu>
- [66] Feng Tian, Flora D. Salim, and Hao Xue. 2025. TradingGroup: A Multi-Agent Trading System with Self-Reflection and Data-Synthesis. arXiv:2508.17565 [cs.AI] <https://arxiv.org/abs/2508.17565>
- [67] Meiyun Wang, Kiyoshi Izumi, and Hiroki Sakaji. 2024. LLMFactor: Extracting Profitable Factors through Prompts for Explainable Stock Movement Prediction. arXiv:2406.10811 [cs.CL] <https://arxiv.org/abs/2406.10811>
- [68] Saizhuo Wang, Hang Yuan, Lionel M. Ni, and Jian Guo. 2024. QuantAgent: Seeking Holy Grail in Trading by Self-Improving Large Language Model. arXiv:2402.03755 [cs.AI] <https://arxiv.org/abs/2402.03755>
- [69] Saizhuo Wang, Hang Yuan, Leon Zhou, Lionel Ni, Heung Yeung Shum, and Jian Guo. 2025. Alpha-GPT: Human-AI Interactive Alpha Mining for Quantitative

- Investment. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. <https://aclanthology.org/2025.emnlp-demos.14/>
- [70] Yiqing Wang, Dehao Dai, Ding Ma, and Kerui Geng. 2026. Macro Economists in the Machine: A Multi-Agent LLM Framework for Commodity-Related ETF Portfolio Construction. *arXiv:2606.08283 [q-fin.PM]* <https://arxiv.org/abs/2606.08283>
- [71] Yanlong Wang, Jian Xu, Hongkang Zhang, Shao-Lun Huang, Danny Dongning Sun, and Xiao-Ping Zhang. 2026. FactorMiner: A Self-Evolving Agent with Skills and Experience Memory for Financial Alpha Discovery. *arXiv:2602.14670 [q-fin.TR]* <https://arxiv.org/abs/2602.14670>
- [72] Yining Wang, Jinman Zhao, and Yuri Lawryshyn. 2024. GPT-Signal: Generative AI for Semi-automated Feature Engineering in the Alpha Research Process. *arXiv:2410.18448 [cs.CE]* <https://arxiv.org/abs/2410.18448>
- [73] Zhangyuhua Weng, Shengli Zhang, Taotao Wang, and Yihan Xia. 2026. AlphaLogics: A Market Logic-Driven Multi-Agent System for Scalable and Interpretable Alpha Factor Generation. *arXiv:2603.20247 [q-fin.CP]* <https://arxiv.org/abs/2603.20247>
- [74] Siyi Wu, Junqiao Wang, Zhaoyang Guan, Leyi Zhao, Xinyuan Song, Xinyu Ying, Dexu Yu, Jinhao Wang, Hanlin Zhang, Michele Pak, Yangfan He, Yi Xin, Jianhui Wang, and Tianyu Shi. 2025. MountainLion: A Multi-Modal LLM-Based Agent System for Interpretable and Adaptive Financial Trading. *arXiv:2507.20474 [q-fin.TR]* <https://arxiv.org/abs/2507.20474>
- [75] Yijia Xiao, Edward Sun, Tong Chen, Fang Wu, Di Luo, and Wei Wang. 2025. Trading-R1: Financial Trading with LLM Reasoning via Reinforcement Learning. *arXiv:2509.11420 [q-fin.TR]* <https://arxiv.org/abs/2509.11420>
- [76] Yijia Xiao, Edward Sun, Di Luo, and Wei Wang. 2024. TradingAgents: Multi-Agents LLM Financial Trading Framework. *arXiv:2412.20138 [q-fin.TR]* <https://arxiv.org/abs/2412.20138>
- [77] Zhuohan Xie, Rania Elbadry, Fan Zhang, Georgi Georgiev, Xueqing Peng, Lingfei Qian, Jimin Huang, Dimitar Dimitrov, Vanshika Jani, Yuyang Dai, Jiahui Geng, Yuxia Wang, Ivan Koychev, Veselin Stoyanov, and Preslav Nakov. 2026. The CLEF-2026 FinMMEval Lab: Multilingual and Multimodal Evaluation of Financial AI Systems. *arXiv:2602.10886 [cs.CL]* <https://arxiv.org/abs/2602.10886>
- [78] Fei Xiong, Xiang Zhang, Aosong Feng, Siqi Sun, and Chenyu You. 2025. QuantHarness: Price-Driven Multi-Agent LLMs for High-Frequency Trading. *arXiv:2509.09995 [cs.CE]* <https://arxiv.org/abs/2509.09995>
- [79] Guojun Xiong, Zhiyang Deng, Keyi Wang, Yupeng Cao, Haohang Li, Yangyang Yu, Xueqing Peng, Mingquan Lin, Kaleb E. Smith, Xiao-Yang Liu, Jimin Huang, Sophia Ananiadou, and Qianqian Xie. 2025. FLAG-TRADER: Fusion LLM-Agent with Gradient-based Reinforcement Learning for Financial Trading. In *Findings of the Association for Computational Linguistics: ACL 2025*. <https://aclanthology.org/2025.findings-acl.716/>
- [80] Feng Xu, Yan Yin, Xinyu Zhang, Tianyuan Liu, Shengyi Jiang, and Zongzhang Zhang. 2024. Alpha²: Discovering Logical Formulaic Alphas using Deep Reinforcement Learning. *arXiv:2406.16505 [q-fin.CP]* <https://arxiv.org/abs/2406.16505>
- [81] Weicheng Xue. 2026. Representation Signatures and Risk-Feedback Alignment in LLM Trading Agents. *arXiv:2605.28850 [cs.LG]* <https://arxiv.org/abs/2605.28850>
- [82] Lewen Yan, Jilin Mei, Tianyi Zhou, Lige Huang, Jie Zhang, Dongrui Liu, and Jing Shao. 2025. TradeTrap: Are LLM-based Trading Agents Truly Reliable and Faithful? *arXiv:2512.02261 [cs.AI]* <https://arxiv.org/abs/2512.02261>
- [83] Xu Yang, Xiao Yang, Shikai Fang, Yifei Zhang, Jian Wang, Bowen Xian, Qizheng Li, Jingyuan Li, Minrui Xu, Yuante Li, Haoran Pan, Yuge Zhang, Weiqing Liu, Yelong Shen, Weizhu Chen, and Jiang Bian. 2025. R&D-Agent: An LLM-Agent Framework Towards Autonomous Data Science. *arXiv:2505.14738 [cs.AI]* <https://arxiv.org/abs/2505.14738>
- [84] Yuxuan Ye, Jun Han, Ao Hu, Juncheng Bu, Yiyi Chen, Liangjian Wen, Danilo Mandic, Danny Dongning Sun, Xu Yinghui, and Zenglin Xu. 2026. The Alpha Illusion: Reported Alpha from LLM Trading Agents Should Not Be Treated as Deployment Evidence. *arXiv:2605.16895 [cs.CE]* <https://arxiv.org/abs/2605.16895>
- [85] Jingyang Yi, Jian Yang, Yifei Jin, Yuqi Li, and Jian Li. 2026. AlphaSchema: Exploring the Space of Trading Semantics for LLM-Based Alpha Mining. *arXiv:2607.26642 [cs.AI]* <https://arxiv.org/abs/2607.26642>
- [86] Hao-fei Yu, Fenghai Li, and Jiaxuan You. 2025. LiveTradeBench: Seeking Real-World Alpha with Large Language Models. *arXiv:2511.03628 [q-fin.TR]* <https://arxiv.org/abs/2511.03628>
- [87] Hang Yu, Zifan Zheng, Jeff Z. Pan, Tongliang Liu, Zhiyong Wang, and Fengxiang He. 2026. AlphaMemo: Structured Search-Process Memory for Self-Evolving Alpha Mining Agents. *arXiv:2606.20625 [cs.AI]* <https://arxiv.org/abs/2606.20625>
- [88] Shuo Yu, Hongyan Xue, Xiang Ao, Feiyang Pan, Jia He, Dandan Tu, and Qing He. 2023. Generating Synergistic Formulaic Alpha Collections via Reinforcement Learning. In *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. doi:10.1145/3580305.3599831
- [89] Yangyang Yu, Haohang Li, Zhi Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Suchow, and Khaldoun Khashanah. 2023. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. *arXiv:2311.13743 [q-fin.CP]* <https://arxiv.org/abs/2311.13743>
- [90] Yangyang Yu, Zhiyuan Yao, Haohang Li, Zhiyang Deng, Yupeng Cao, Zhi Chen, Jordan W. Suchow, Rong Liu, Zhenyu Cui, Zhaozhao Xu, Denghui Zhang, Koduvayur Subbalakshmi, Guojun Xiong, Yueru He, Jimin Huang, Dong Li, and Qianqian Xie. 2024. FinCon: A Synthesized LLM Multi-Agent System with Conceptual Verbal Reinforcement for Enhanced Financial Decision Making. *arXiv:2407.06567 [cs.CL]* <https://arxiv.org/abs/2407.06567>
- [91] Hang Yuan, Saizhuo Wang, and Jian Guo. 2024. Alpha-GPT 2.0: Human-in-the-Loop AI for Quantitative Investment. *arXiv:2402.09746 [q-fin.CP]* <https://arxiv.org/abs/2402.09746>
- [92] Yishuo Yuan, Jiayi Sheng, Sirui Zeng, Jiaqi Wang, and Jiaheng Liu. 2026. AlphaCrafter: Harnessing Multi-Agent Workflows for Cross-Sectional Quantitative Trading. *arXiv:2605.05580 [cs.AI]* <https://arxiv.org/abs/2605.05580>
- [93] Lingzhe Zhang, Tong Jia, Yunpeng Zhai, Zixuan Xie, Chiming Duan, Minghua He, Philip S. Yu, and Ying Li. 2026. From Feedback Loops to Policy Updates: Reinforcement Fine-Tuning for LLM-Based Alpha Factor Discovery. *arXiv:2605.15412 [cs.CE]* <https://arxiv.org/abs/2605.15412>
- [94] Tianping Zhang, Yuanqi Li, Yifei Jin, and Jian Li. 2020. AutoAlpha: an Efficient Hierarchical Evolutionary Algorithm for Mining Alpha Factors in Quantitative Investment. *arXiv:2002.08245 [q-fin.CP]* <https://arxiv.org/abs/2002.08245>
- [95] Wentao Zhang, Lingxuan Zhao, Haochong Xia, Shuo Sun, Jiaze Sun, Molei Qin, Xinyi Li, Yuqing Zhao, Yilei Zhao, Xinyu Cai, Longtao Zheng, Xinrun Wang, and Bo An. 2024. A Multimodal Foundation Agent for Financial Trading: Tool-Augmented, Diversified, and Generalist. *arXiv:2402.18485 [q-fin.TR]* <https://arxiv.org/abs/2402.18485>
- [96] Wentao Zhang, Mingxuan Zhao, Jincheng Gao, Jieshun You, Huaiyu Jia, Yilei Zhao, Bo An, and Shuo Sun. 2026. AlphaForgeBench: Benchmarking End-to-End Trading Strategy Design with Large Language Models. *arXiv:2602.18481 [q-fin.TR]* <https://arxiv.org/abs/2602.18481>
- [97] Tianjiao Zhao, Jingrao Lyu, Stokes Jones, Harrison Garber, Stefano Pasquali, and Dhagash Mehta. 2025. AlphaAgents: Large Language Model based Multi-Agents for Equity Portfolio Constructions. *arXiv:2508.11152 [q-fin.ST]* <https://arxiv.org/abs/2508.11152>
- [98] Yuxuan Zhao, Sijia Chen, and Ningxin Su. 2026. PortBench: A Correlation-Aware, Full-Pipeline Benchmark for LLM-Driven Portfolio Management. *arXiv:2605.27887 [cs.AI]* <https://arxiv.org/abs/2605.27887>
- [99] Taojie Zhu, Wentao Zhao, Rui Sun, Beidi Luan, Jiacheng Lu, Sinuo Wang, Jing Li, Daxin Jiang, Yonghong He, and Zuo Bai. 2026. From Knowing to Doing: A Memory-Controlled Benchmark for LLM Trading Agents on Stock Markets. *arXiv:2605.28359 [cs.AI]* <https://arxiv.org/abs/2605.28359>
- [100] Mingxi Zou, Jiaxiang Chen, Aotian Luo, Jingyi Dai, Chi Zhang, Dongning Sun, and Zenglin Xu. 2026. FinEvo: From Isolated Backtests to Ecological Market Games for Multi-Agent Financial Strategy Evolution. *arXiv:2602.00948 [physics.soc-ph]* <https://arxiv.org/abs/2602.00948>